

Get Familiar with FDB Source Code

Please complete the following sub-tasks (the order is just a recommendation, you could complete them in any order):

1. Download the FDB source code at <https://github.com/apple/foundationdb> and build it locally in your machine (either Linux or Mac; If you have Mac M1/M2, this sub-task is even more challenging). Prepare a document that lists the steps you did and the commands you executed.
2. From the source code built at Sub-task 1, launch the FDB server and use the fdbcli to connect to it. Issue simple KV operations (get, set, getrange, delete) via the fdbcli tool.
3. Use the FDB C library and write a C program (named the file basic_ops.c) to perform basic operations (get/set/getrange/delete) to the FDB server. The document for FDB C API could be found at: <https://apple.github.io/foundationdb/api-c.html>
4. Measure single getrange performance:
 - a. Create a file single_get_range.c and use it for this sub-task
 - b. Store 10k key-value pairs (key_i, val_i) in FDB
 - c. Retrieve all 10k key-value pairs by executing a getrange on \x00\xff
 - d. Modify getrange to use different modes (WANT_ALL, EXACT, ITERATOR, etc..) and report the response time of each execution
5. Compare single getrange vs multiple getranges sent in parallel:
 - a. Create a file single_vs_multi_ranges.c and use it for this sub-task
 - b. With the 10k key-value pairs stored in sub-task 5, define 10 ranges R1, R2, ..., R10 such that executing getrange on each of those 10 ranges returns exactly 1k key-value pairs.
 - c. Execute getRange() requests on these 10 ranges in parallel
 - d. Repeat the execution with different streaming modes (WANT_ALL, EXACT, ITERATOR, etc..) and report the response time of each execution. How are the numbers compared with the case in sub-task 5 where we only issue one getrange on \x00\xff?
6. Understand read snapshot:
 - a. Create a class read_snapshot.c and use it for this sub-task
 - b. Store any arbitrary 4 key-value pairs in the database. Denote the keys as K1, K2, K3 and K4
 - c. Start a transaction T1 and read several keys (let's say K1, K2, K3)
 - d. On another thread, start a transaction T4 that updates the values of K2, K4
 - e. Commit transaction T1. Would T1 be aborted? Explain why.
 - f. Commit transaction T2. Would T2 be aborted? Explain why.
7. Understand transaction conflict:
 - a. Create a class tx_conflict.c and use it for this sub-task
 - b. Store any arbitrary 2 key-value pairs in the database. Denote the keys as K1 and K2.
 - c. Start a transaction T1 to read K1 and update the value of K2
 - d. Start a transaction T2 to read K2 and update the value of K1
 - e. Commit T2. Would T2 be aborted? Explain why.
 - f. Commit T1. Would T1 be aborted? Explain why.

Note:

- If you cannot complete sub-task 1, please install FDB server and client directly from here <https://apple.github.io/foundationdb/client-design.html>
- FDB stream modes: <https://apple.github.io/foundationdb/api-c.html#c.FDBStreamingMode>

Delivery:

- A detailed document for sub-task 1, 2 and report numbers for sub-task 4-7.
- Source code of the other sub-tasks (3, 4, 5, 6 and 7)
 - basic_ops.c
 - single_get_range.c
 - single_vs_multi_ranges.c
 - read_snapshot.c
 - tx_conflict.c